

HomeWork 6

HomeWork 6

ISYE 6501

Question 9.1

Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2. You can use the R function `prcomp` for PCA. (Note that to first scale the data, you can include `scale. = TRUE` to scale as part of the PCA function. Don't forget that, to make a prediction for the new city, you'll need to unscale the coefficients (i.e., do the scaling calculation in reverse)!))

```
library(kernlab)
library(kknn)
library(lattice)
library(ggplot2)
library(outliers)
library(dplyr)

library(leaps)
```

Next we will load the data and look at the data structure.

[illegible]

```
## 3 24.3006 578
## 4 29.9012 1969
## 5 21.2998 1234
## 6 20.9995 682
```

Lets examine the data for correlations using a visualisation.


```
library (GGally)

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

##
## Attaching package: 'GGally'

## The following object is masked from 'package:dplyr':
##
##   nasa

ggpairs(usrime, columns = c("Po1", "Po2", "U1", "Ineq"))
```



PCA is used when we have large number of features in the hundreds or thousands, and we cannot make a reasoable model with the limited number of data points we have. In the current data set PCA would be over kill and will probably lead to over fitting. PCA also has the added advantage of reducing correlation by giving higher priority to more relavent features. PCA doesnt perform well with binary features

```
usrime$So <- NULL
#head(usrime)
```

```
pca.out = prcomp(uscrime[,1:14],scale = TRUE) # principal component analysis model
pca.out$sdev # gives standard deviations


## [1] 2.32616392 1.65127377 1.41578394 1.03669703 0.96745274 0.74049025
## [7] 0.56414623 0.54675066 0.44751244 0.42747346 0.35944524 0.31852077
## [13] 0.25159368 0.06802123

variance <- pca.out$sdev^2 # get back eigenvalues

summary(pca.out)

## Importance of components:
##
##          PC1      PC2      PC3      PC4      PC5      PC6      PC
7
## Standard deviation    2.3262 1.6513 1.4158 1.03670 0.96745 0.74049 0.5641
5
## Proportion of Variance 0.3865 0.1948 0.1432 0.07677 0.06685 0.03917 0.0227
3
## Cumulative Proportion 0.3865 0.5813 0.7244 0.80121 0.86806 0.90723 0.9299
6
##          PC8      PC9      PC10     PC11     PC12     PC13     P
C14
## Standard deviation    0.54675 0.4475 0.42747 0.35945 0.31852 0.25159 0.06
802
## Proportion of Variance 0.02135 0.0143 0.01305 0.00923 0.00725 0.00452 0.00
033
## Cumulative Proportion 0.95132 0.9656 0.97867 0.98790 0.99515 0.99967 1.00
000

screeplot(pca.out) # Scree plot shows variance explained per principal
```



From the graph and the summary looks like the first 5 Principal components have considerably higher variance. We will use these 4 PCs in our regression model.

```

#Top 5 principal components
pca1 <- pca.out$x[,1:5]
#pca1

uscrimePC <- cbind(pca1, uscrime["Crime"])
#head(uscrimePC)
#Running a linear model
modelPCA <- lm(Crime~.,data = as.data.frame(uscrimePC))
summary(modelPCA)

##
## Call:
## lm(formula = Crime ~ ., data = as.data.frame(uscrimePC))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -439.96 -181.93   3.13  177.53  444.64
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  905.085     36.364   24.890 < 2e-16 ***
## PC1           76.750     15.802    4.857 1.77e-05 ***
## PC2          -57.648     22.260   -2.590  0.0132 *
## PC3           24.313     25.962    0.936  0.3545
## PC4           3.786     35.456    0.107  0.9155
## PC5          -235.831     37.994   -6.207 2.20e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 249.3 on 41 degrees of freedom
## Multiple R-squared:  0.6297, Adjusted R-squared:  0.5845
## F-statistic: 13.94 on 5 and 41 DF, p-value: 5.685e-08

```

In the above model our adjusted R square is abysmal. Lets do some feature selection on the Principal components. We will use `ols_subset` to Select the subset of predictors have the largest R2 value or the smallest mean squared error.

```

#Get all principal components from pca output
pc_full <- pca.out$x[,]
head(pc_full)

##              PC1              PC2              PC3              PC4              PC5              PC6
## [1,] -3.893446 -1.29197714 -1.10138991  0.85371781 -0.2582188 -0.2220763
## [2,]  0.971018  0.69084709 -0.05783388 -0.36390197 -1.1311282  0.6335173
## [3,] -3.946974  0.08584861 -0.36356482  0.58382631  0.3690864 -0.7120597
## [4,]  3.951699 -2.29488126  0.22720214 -0.04741697 -1.7050714 -0.6698441
## [5,]  1.647039  1.44338543  1.25954528  0.64890563 -0.1162943  0.4327958
## [6,]  2.861367 -0.11765453  0.51823342  1.44813377  1.0732058  0.6906828
##              PC7              PC8              PC9              PC10              PC11              PC1
## [1,]  0.48212291  0.07191639 -0.46939569 -0.05399584 -0.2035171  0.0358941
## [2,] -0.20182701 -0.37733855 -0.24477852  0.02972945 -0.3578821 -0.1584916
## [3,]  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111
## [4,]  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111
## [5,]  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111
## [6,]  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111  0.11111111

```

```
## [3,] 0.01482101 -0.37633010 0.08792458 -0.40993363 0.1155926 0.3097035
2
## [4,] -0.74188485 1.42942026 -0.20222448 -0.01001613 -0.2187344 -0.1864198
0
## [5,] -0.28482158 0.38994226 -0.49201684 -0.22200390 -0.3052962 0.6775270
4
## [6,] -0.06932221 -0.70963951 -0.13164159 0.12061204 0.7156631 -0.0485135
4
```

Now that we have our PC values we need to reverse engineer to the coefficients to get the right features. Below is a work through of the math.

using a toy data set.

$$\frac{z}{\sigma_x \sigma_y} = \frac{a_x}{\sigma_x} \left(\frac{x - \mu_x}{\sigma_x} \right) + \frac{a_y}{\sigma_y} \left(\frac{y - \mu_y}{\sigma_y} \right) + \frac{I + \mu_x + \mu_y}{\sigma_x \sigma_y}$$

where a_x and a_y are x and y coefficients
 I is the intercept.

$$z = a_x \sigma_x x_s + a_y \sigma_y y_s + I + \mu_x + \mu_y$$

where x_s and y_s are scaled data points

Now to apply the transformation where the rotation matrix is $\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$

$$\begin{aligned} x' &= x_s \cos \theta - y_s \sin \theta \\ y' &= x_s \sin \theta + y_s \cos \theta \end{aligned}$$

inversely $x_s = \cos \theta x' - \sin \theta y'$
 $y_s = \sin \theta x' + \cos \theta y'$ here x' and y' are principle components

$$z = a_x \sigma_x (\cos \theta x' - \sin \theta y') + a_y \sigma_y (\sin \theta x' + \cos \theta y') + I + \mu_x + \mu_y$$

Say from our Linear regression we get:

$$z = a'_x x' + a'_y y' + I'$$

where a'_x, a'_y and I' are based on PC's.

$$a_x \sigma_x \cos \theta + a_y \sigma_y \sin \theta = a'_x \quad \text{and} \quad -a_x \sigma_x \sin \theta + a_y \sigma_y \cos \theta = a'_y$$

To rotate back and get a_x and a_y

$$\begin{bmatrix} \sigma_x a_x \\ \sigma_y a_y \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} a'_x \\ a'_y \end{bmatrix}$$

to unscale the coefficients
 divide the result by $sd(\sigma)$

In order to get the scaled coefficients in original features we need to multiply the coefficient vector to the rotation matrix of the PCA. To unscale the coefficient matrix we need to divide by standard deviation. For the intercept we only divide by the PC intercept by the sum of the means.

```
# Rotation vector for the first 5 PCs
rotation_vec <- pca.out$rotation[,1:5]
```

```

# Coefficients of PCA model without the Intercept
PCA_coef <- modelPCA$coefficients[2:6]

scaled_coef <- PCA_coef %*% t(rotation_vec)
scaled_coef

##           M           Ed           Po1           Po2           LF           M.F           Pop
NW
## [1,] 64.14904 14.09834 116.4574 110.4209 76.17019 103.1387 63.33765 104.15
81
##           U1           U2    Wealth           Ineq           Prob           Time
## [1,] 3.83432 32.23305 29.57314 25.60389 -27.59591 26.85164

```

Now to get original coefficients and intercepts

```

intercept <- modelPCA$coefficients[1] - sum(scaled_coef * sapply(uscrime[,1:14], mean) / pca.out$scale) #sum(sapply(uscrime[,1:15], mean))
intercept

## (Intercept)
## -5726.224

unscaled_coef <- scaled_coef / pca.out$scale
unscaled_coef

##           M           Ed           Po1           Po2           LF           M.F           Pop           N
W
## [1,] 51.04305 12.60243 39.18621 39.49058 1884.85 35.00099 1.663664 10.1292
7
##           U1           U2    Wealth           Ineq           Prob           Time
## [1,] 212.6777 38.16618 0.03064862 6.417648 -1213.702 3.788914

matrix_data <- as.matrix(uscrime[,1:14])

# estimate is of the form estimate = aX + b
# where a are the coefficients and b is the intercept
estimates <- matrix_data %*% t(unscaled_coef) + intercept
SSE = sum((estimates - uscrime[,15])^2)
SStot = sum((uscrime[,15] - mean(uscrime[,15]))^2)
R2 <- 1 - SSE/SStot
R2

## [1] 0.6296769

# Create the test datapoint manually using the data.frame() function for the
new city
test_point <- data.frame(M = 14.0, Ed = 10.0 , Po1 = 12.0 , Po2 = 15.5, LF = 0.64
0 , M.F = 94.0 , Pop = 150 , NW = 1.1 , U1= 0.120 , U2 = 3.6 , Wealth = 3200, Ineq =
20.1, Prob = 0.04 , Time = 39.0)

# Use the intercepts and coefficients to make a prediction of Crime in the new
city
matrix_test <- as.matrix(test_point)

```

```
prediction = matrix_test %*% t(unscaled_coef) + intercept
prediction

##           [,1]
## [1,] 1443.039
```

As we already know the crime for the new city from the previous Homework, our model is overestimating. As mentioned earlier PCA is ideally suited for large number of data points. With our 50 data points the model is over fitting. Compared to my previous model from 8.2 which gave a value of 728, this model is less accurate. The final prediction for test city using PCA is 1443 with an R-square of 0.629.

```
#Sanity check to see if our reverse PCA calculation was correct
#project new data onto PCA space and run model

pca_test <- scale(test_point, pca.out$center, pca.out$scale) %*% pca.out$rotation
pca_test

pca_test <- as.data.frame(pca_test)
pca_test

##           PC1           PC2           PC3           PC4           PC5           PC6           PC7
## [1,]  1.161658 -2.841351  0.5694485 -1.04263 -1.166522 -2.191452 -0.4660637
##           PC8           PC9          PC10          PC11          PC12          PC13          PC14
## [1,]  0.9344984  0.227878  0.555688 -1.088542  3.504094  0.6951131  1.269102

predict(modelPCA, newdata=pca_test, type="response")

##           1
## 1443.039
```

We get the same response by projecting the data onto the PCA axis.